


Metadata security

Mihai Ordean
Designing and Managing Secure Systems
University of Birmingham


Overview

- Device security
 - Is code on the device vulnerable to exploits ? (e.g. buffer overflows)
 - Is the code authenticated ? (i.e. has not been tampered with)
- Local data security
 - Is the stored data is accessible to everyone? (e.g. encrypted)
 - Is the stored data authenticated?
- Cloud data security
 - How is data stored in the cloud?
 - Who has access to data stored in the cloud?
- **Metadata security**
 - What does metadata reveal about stored data?
 - Can we tamper the metadata?


What can we encrypt?




What can we encrypt?



What can we encrypt?



What can we encrypt?



Protecting dynamic data




Protecting dynamic data




Protecting dynamic data




Protecting static data




Protecting static data




Protecting static data



Protecting static data




Protecting static data



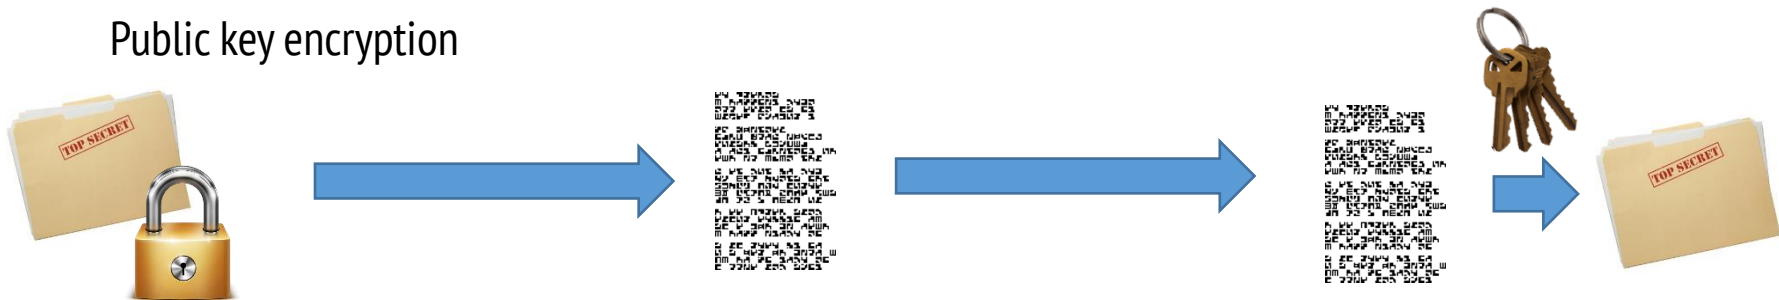
Protecting data from the cloud



Symmetric key encryption

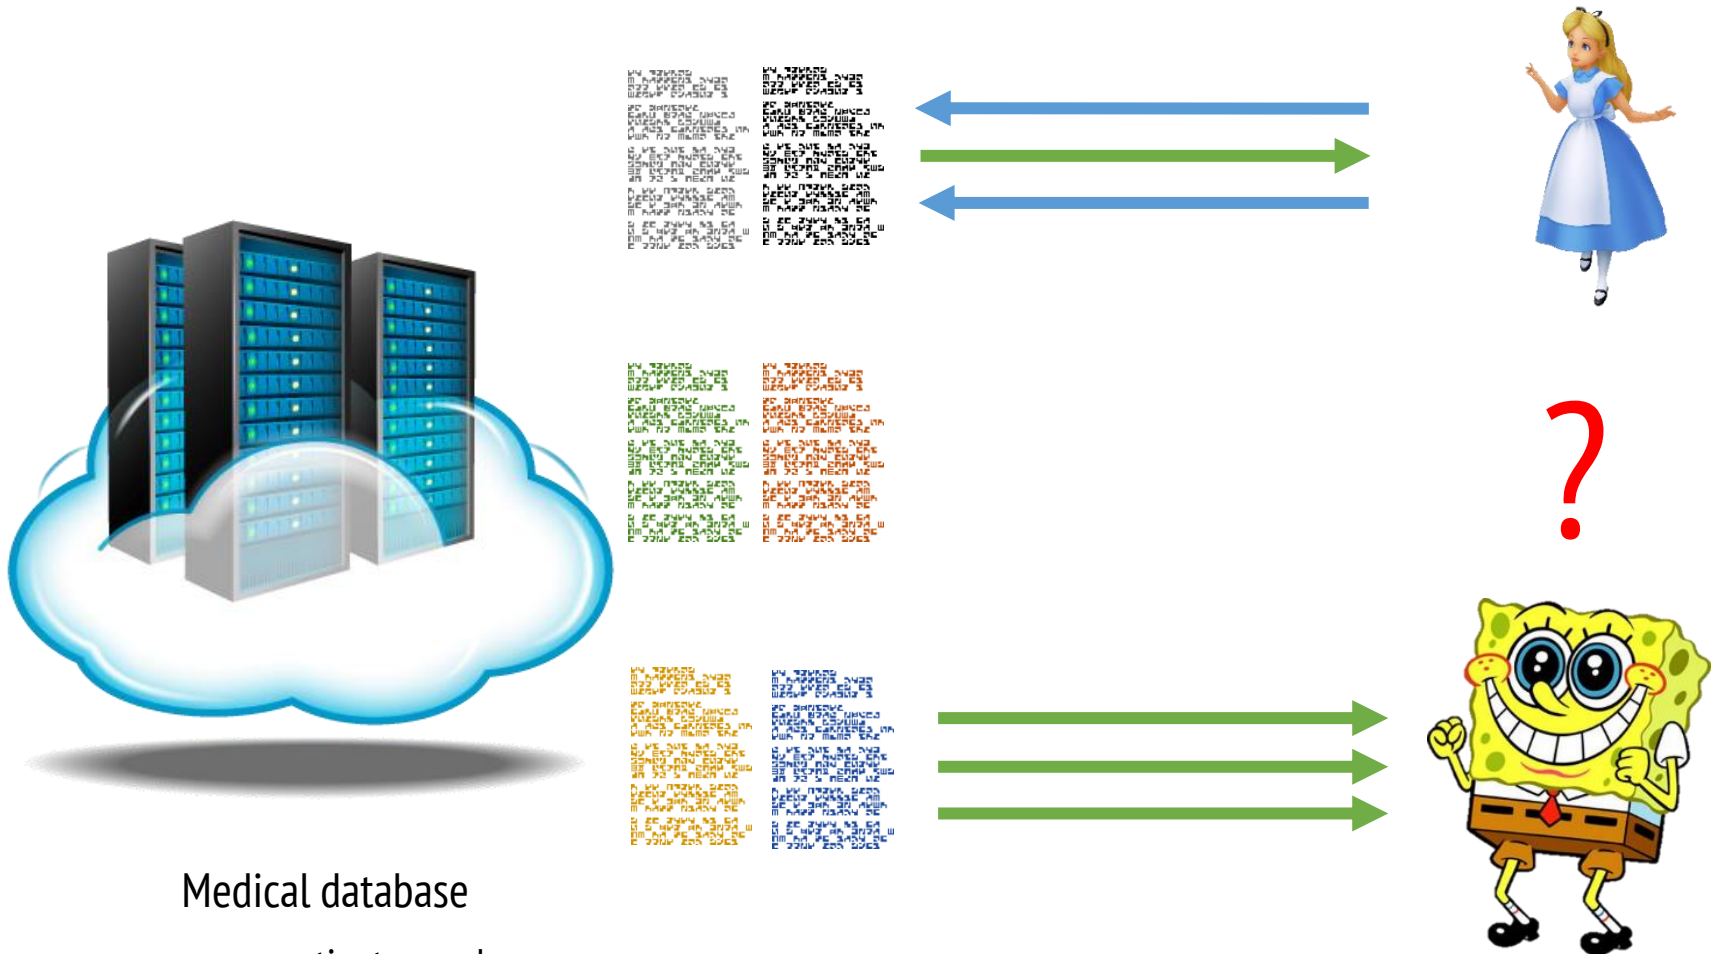


Public key encryption



Is encrypting data enough?


Analysing data access: who is the doctor?



Medical database

- patient records
- insurance records
- appointments

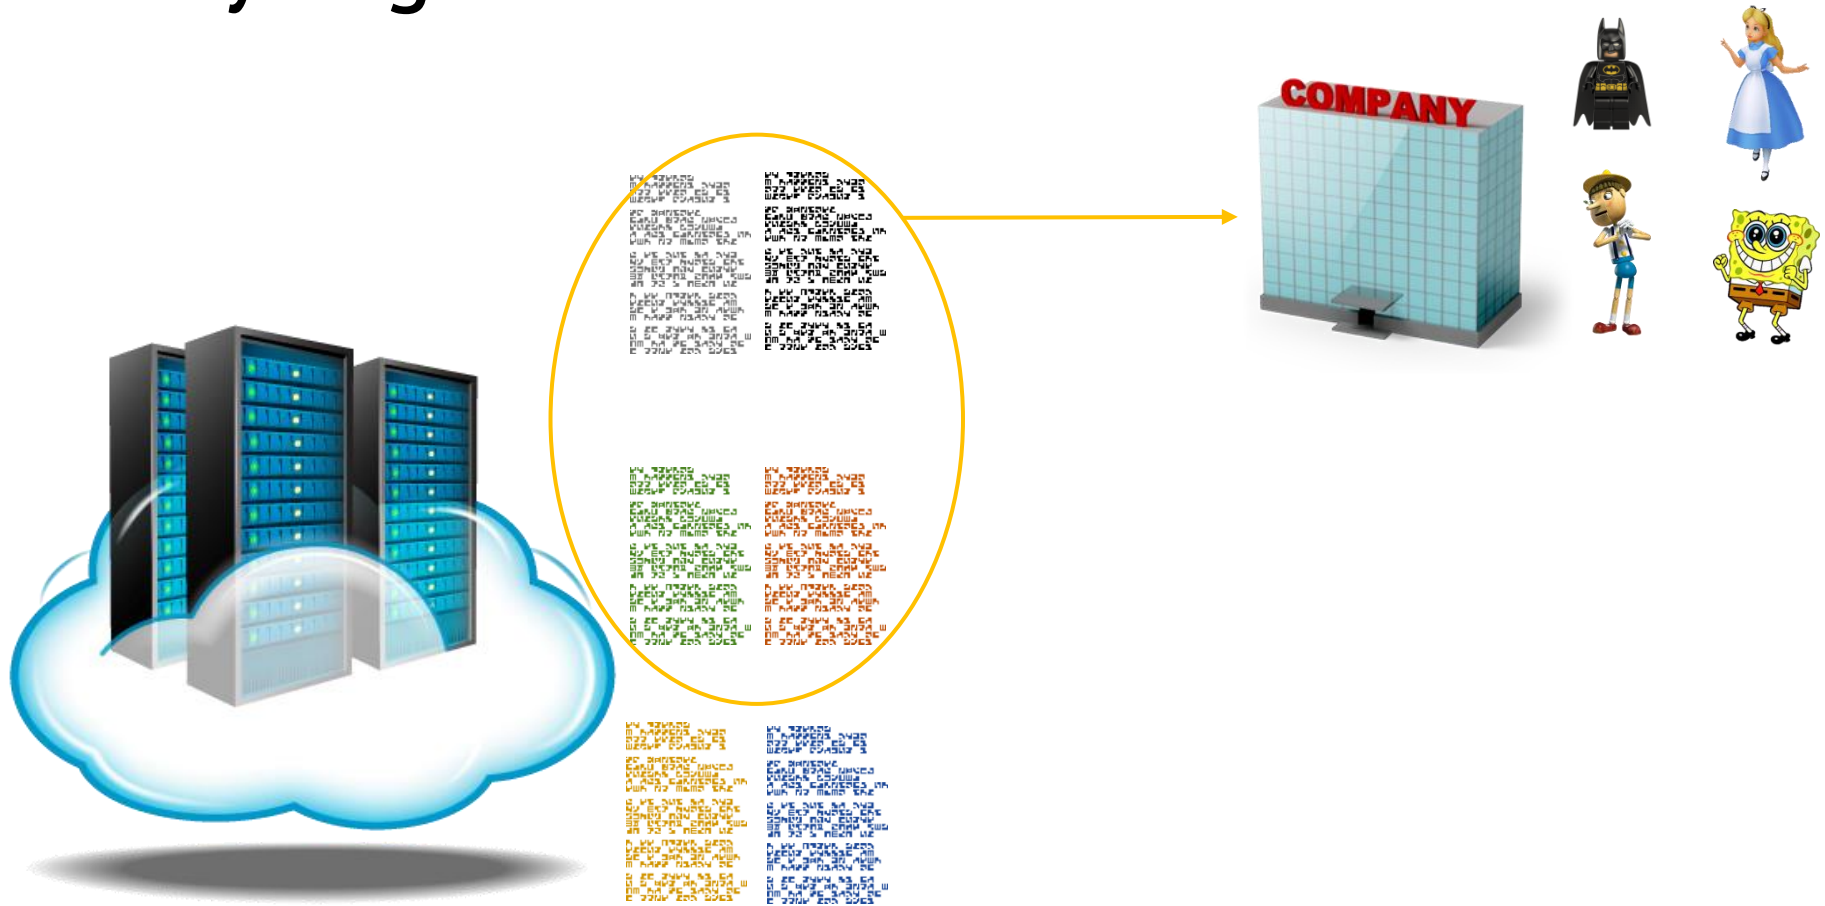
Analysing data access: who owns this ciphertext?



Medical database

- patient records
- insurance records
- appointments

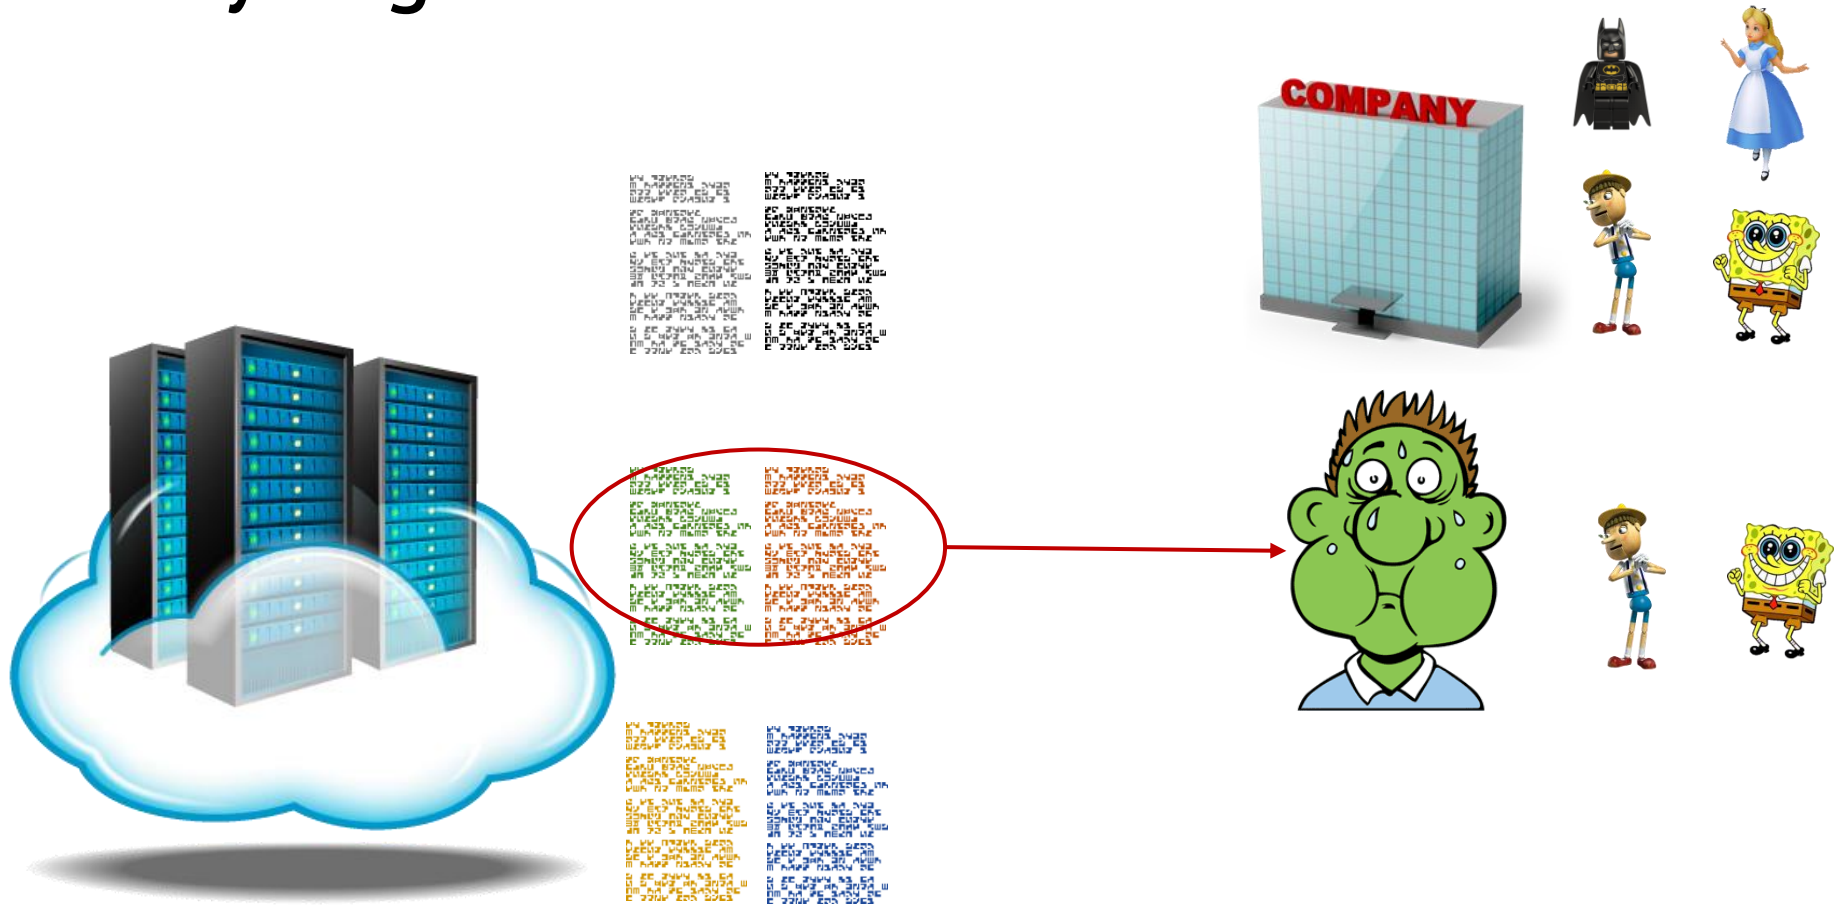
Analysing data access



Medical database

- patient records
- insurance records
- appointments

Analysing data access



Medical database



- patient records
- insurance records
- appointments

Analysing data access

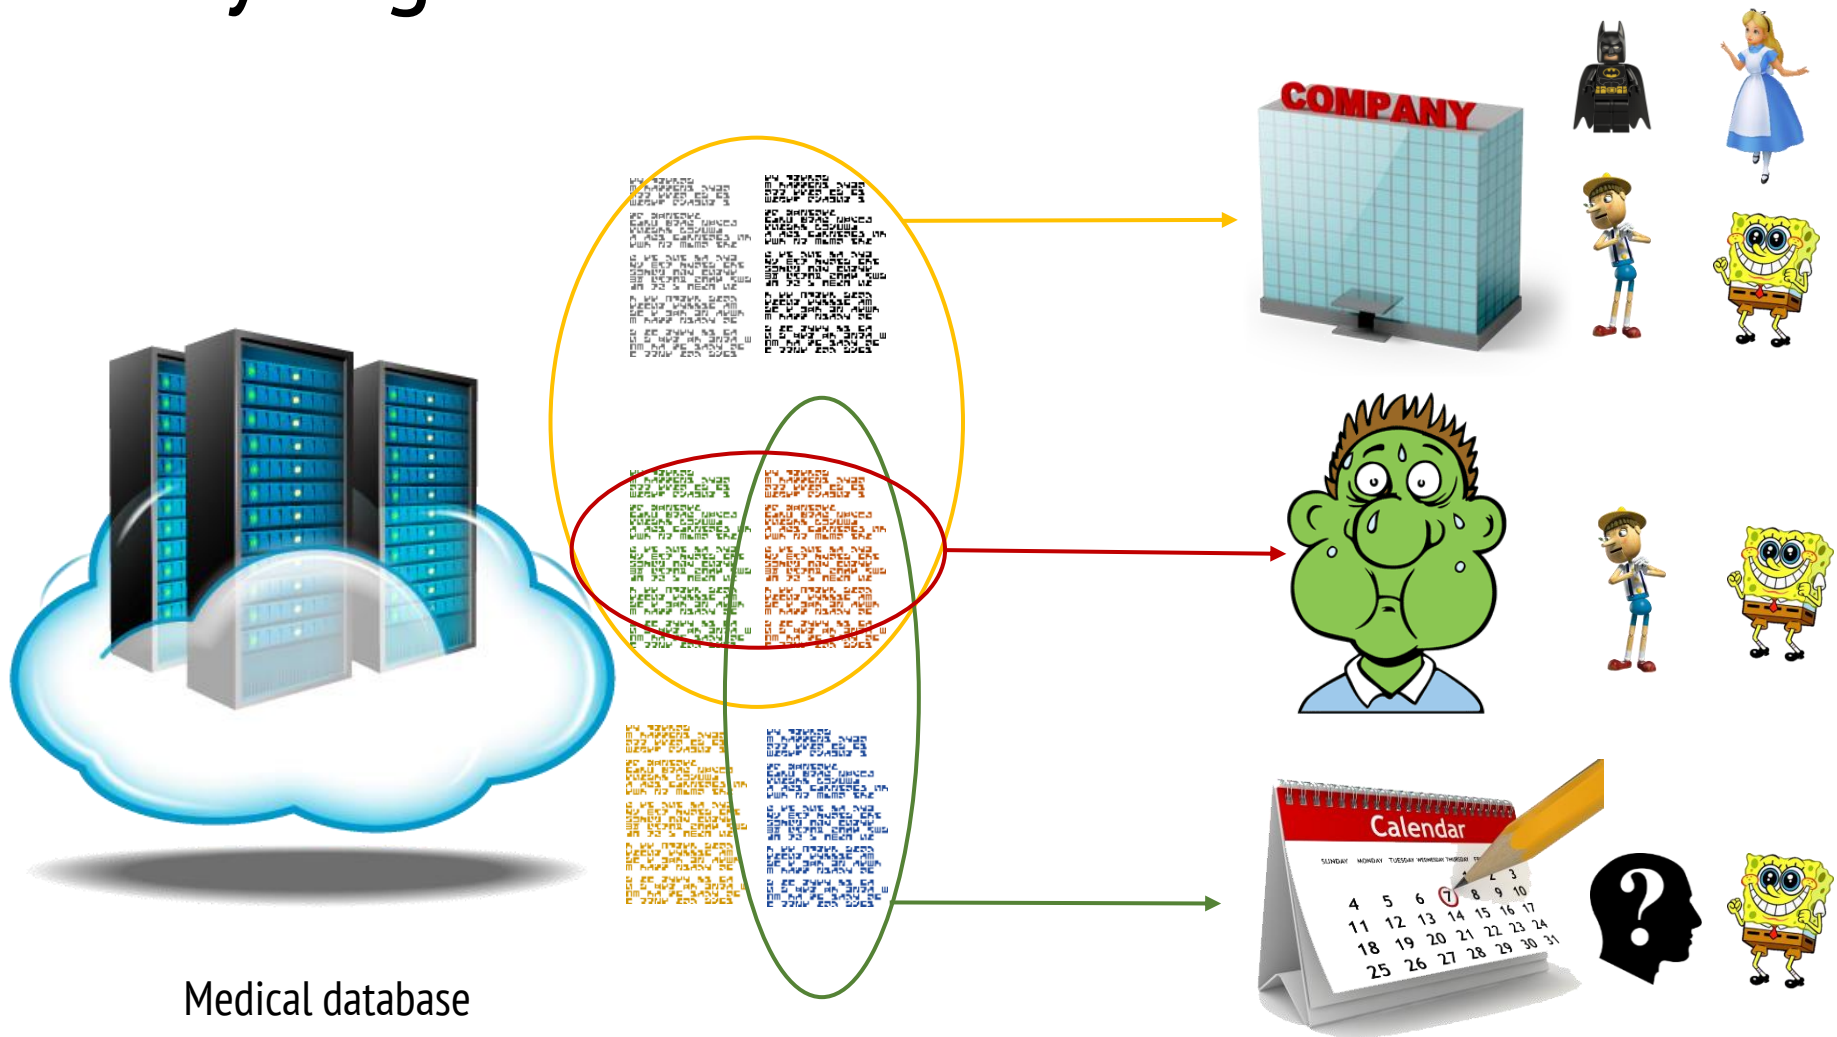


Medical database

- patient records
- insurance records
- appointments




Analysing data access



Medical database

- patient records
- insurance records
- appointments

Analysing data access

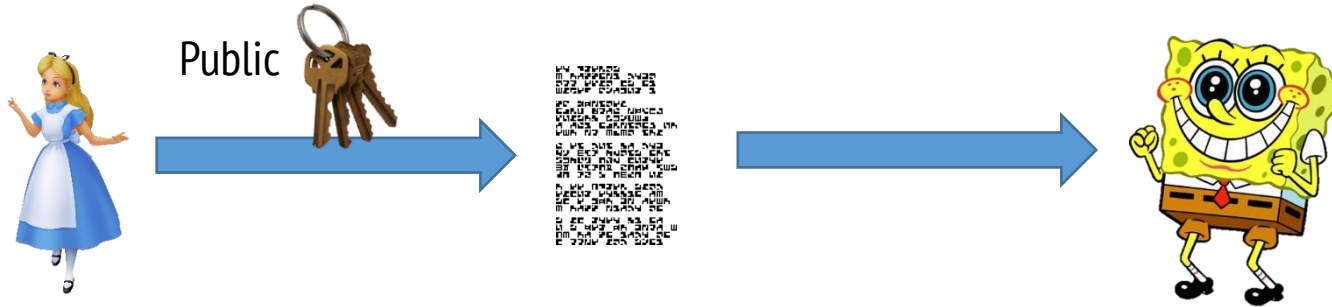


Medical database

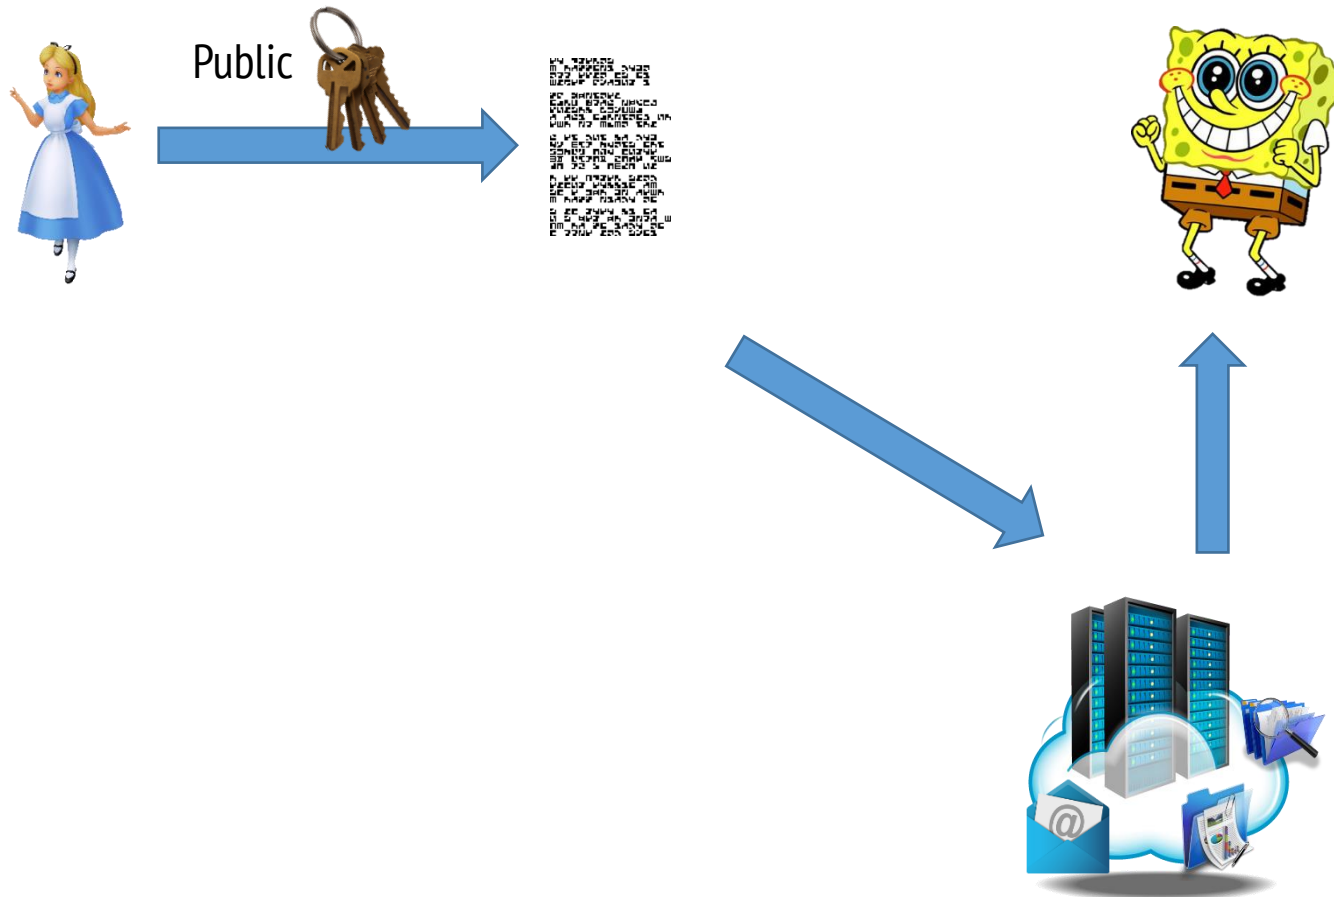
- patient records
- insurance records
- appointments

Cryptography issues

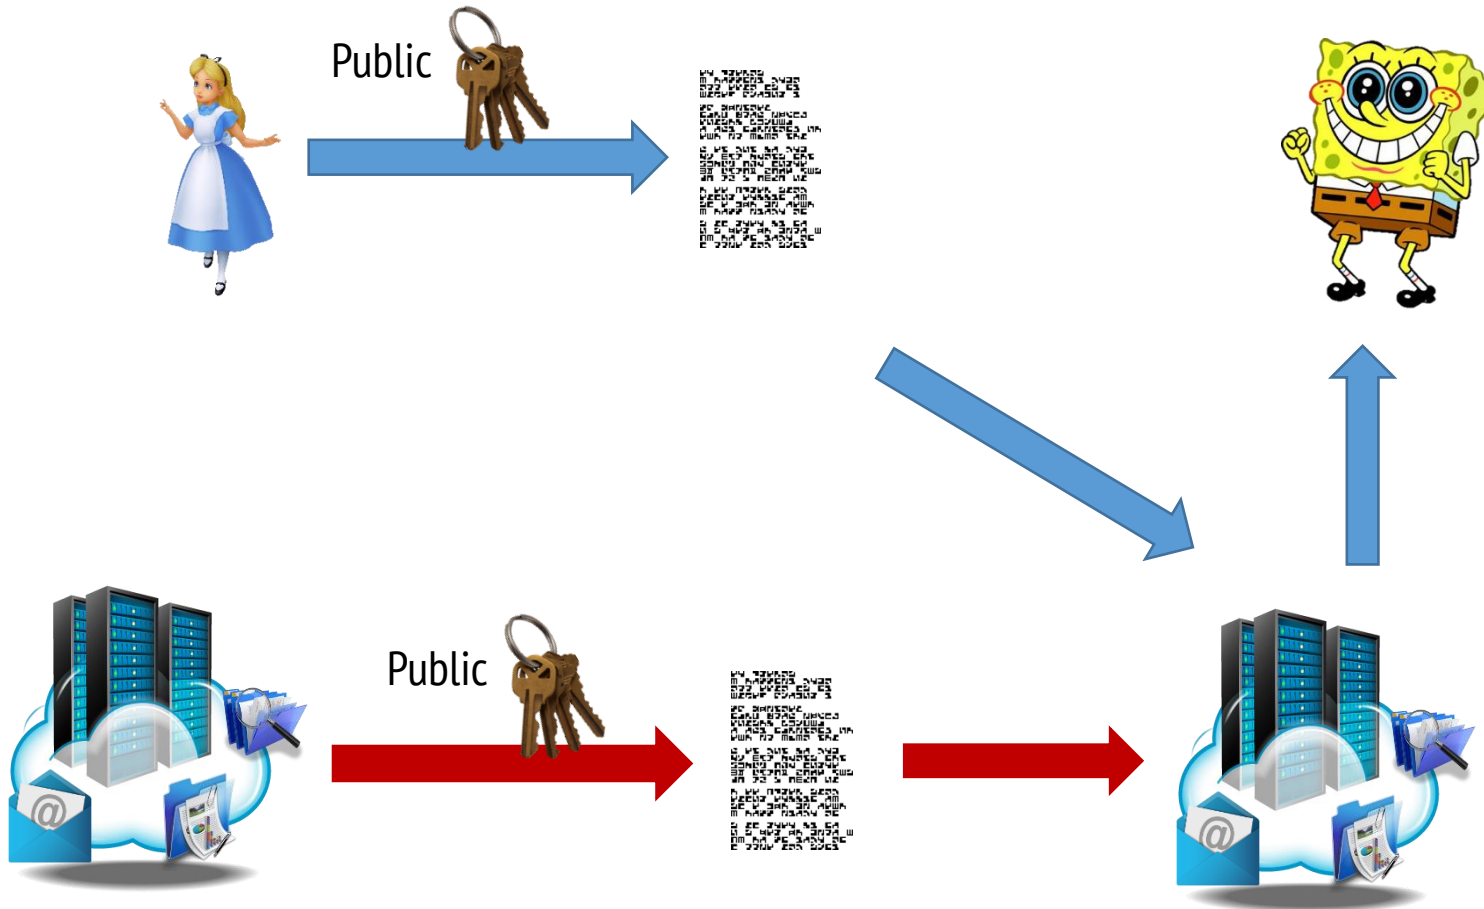
Anyone can use a public key to encrypt



Cryptography issues



Cryptography issues



Cryptography issues

- When protecting metadata, using public key crypto gives you a larger surface of attack.
- Symmetric crypto doesn't have this problem **and** is more efficient.
- Symmetric keys are difficult to share.
- Design schemes based on **symmetric keys** and use **simple public key exchange protocols** to share them.

Just using encryption is not enough

- ✓ **Content security** – the data is encrypted
- ✓ **Metadata security** – ownership information, timestamps, access rights, ciphertext length, etc.
- ✓ **Access pattern security** – when is the data accessed, who accesses the data, how is the data accessed, etc.

Searchable encryption

The challenge (in general)

- Assume we're using Gmail to communicate (with a browser).
- Assume we're using PGP to encrypt email (in browser).
- Can we decrypt email on the fly?
- Can we search through our emails?
- Who performs the search? Is it optimal?

The challenge (in general)

- Assume we're using Gmail to communicate (with a browser).
- Assume we're using PGP to encrypt email (in browser).
- Can we decrypt email on the fly?
 - YES
- Can we search through our emails?
 - Just the ones we decrypted
- Who performs the search? Is it optimal?
 - The client, in browser. Searching on the server would be the optimal choice

Solution?

- Assume we're using Gmail to communicate (with a browser).
- Assume we're using PGP to encrypt email (in browser).
- Generate a searchable index
- Store the index encrypted in the cloud

Solution?

- Assume we're using Gmail to communicate (with a browser).
- Assume we're using PGP to encrypt email (in browser).
- Generate a searchable index
- Store the index encrypted in the cloud
- Client has to download index every time
- Client still does the search, but it's much faster and can be done over all emails.

Solution?

- Assume we're using Gmail to communicate (with a browser).
- Assume we're using PGP to encrypt email (in browser).
- Generate a searchable index
- Store the index encrypted in the cloud
- ~~Client has to download index every time~~
- ~~Client still does the search, but it's much faster and can be done over all emails.~~
- THE SERVER SHOULD DO THE SEARCH! (no download, no computational effort)

Searchable Encryption


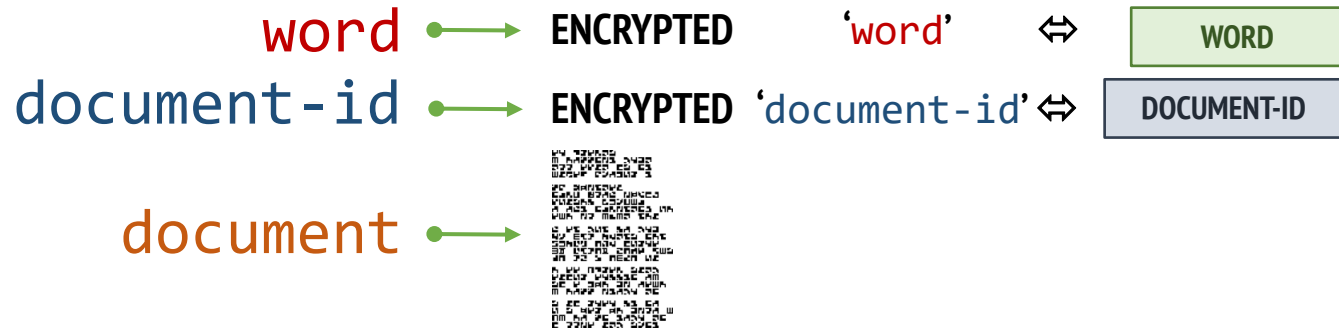


Searching



```
For each document in the database:  
  For each word in document:  
    if word = 'top-secret'  
  
print document-id
```

Encrypting databases

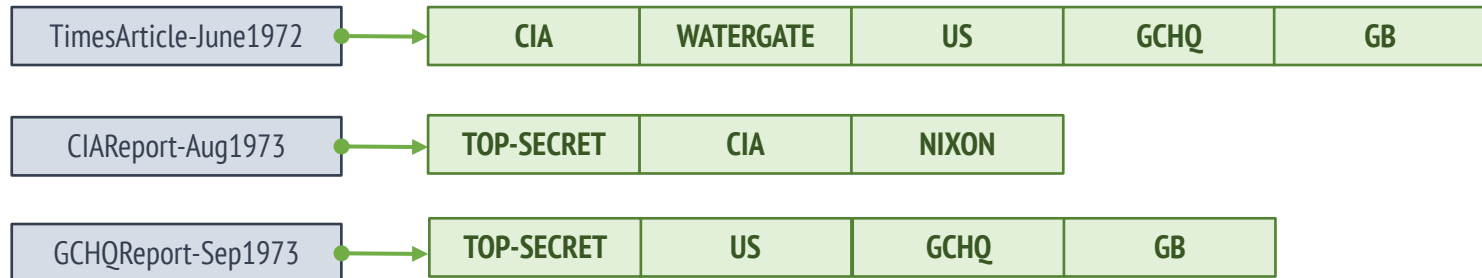


Searchable Encryption

KEYWORDS:

TOP-SECRET	CIA	WATERGATE	NIXON	US	GCHQ	GB
------------	-----	-----------	-------	----	------	----

Forward index



Efficiency of the index

Number of **documents** increases => **time** increases

Number of **keywords** increases => **time** increases

Searchable Encryption

KEYWORDS:

TOP-SECRET	CIA	WATERGATE	NIXON	US	GCHQ	GB
------------	-----	-----------	-------	----	------	----

Inverted index

TOP-SECRET	→	CIAResult-Aug1973	GCHQResult-Sep1973
CIA	→	TimesArticle-June1972	CIAResult-Aug1973
WATERGATE	→	TimesArticle-June1972	
NIXON	→	CIAResult-Aug1973	
US	→	TimesArticle-June1972	GCHQResult-Sep1973
GCHQ	→	TimesArticle-June1972	GCHQResult-Sep1973
GB	→	TimesArticle-June1972	GCHQResult-Sep1973

Efficiency of the index

Number of **keywords** increases => **time** increases

What do we want to protect?

What we search for

KEYWORDS:

TOP-SECRET	CIA	WATERGATE	...
------------	-----	-----------	-----

What is the result of the search query

DOCUMENT NAMES:

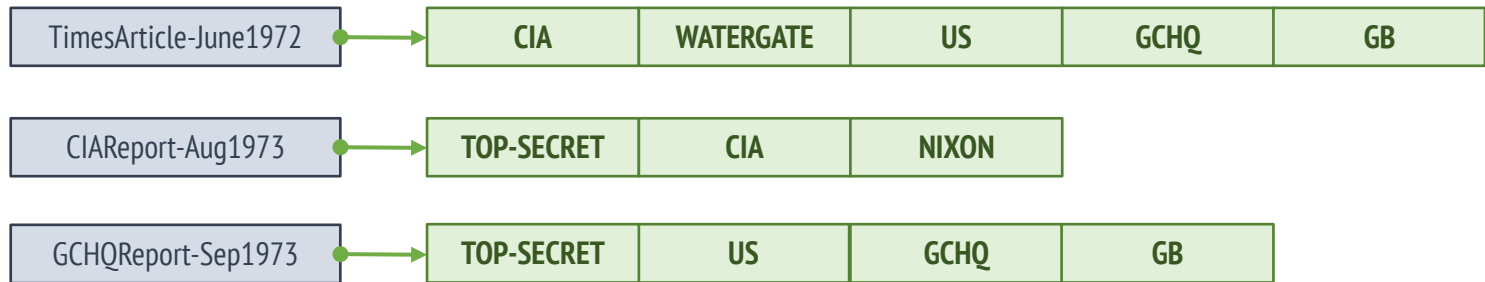
CIAReport-Aug1973	GCHQReport-Sep1973
-------------------	--------------------

How often we search for something

1:	TOP-SECRET
2:	CIA
	...
n:	TOP-SECRET

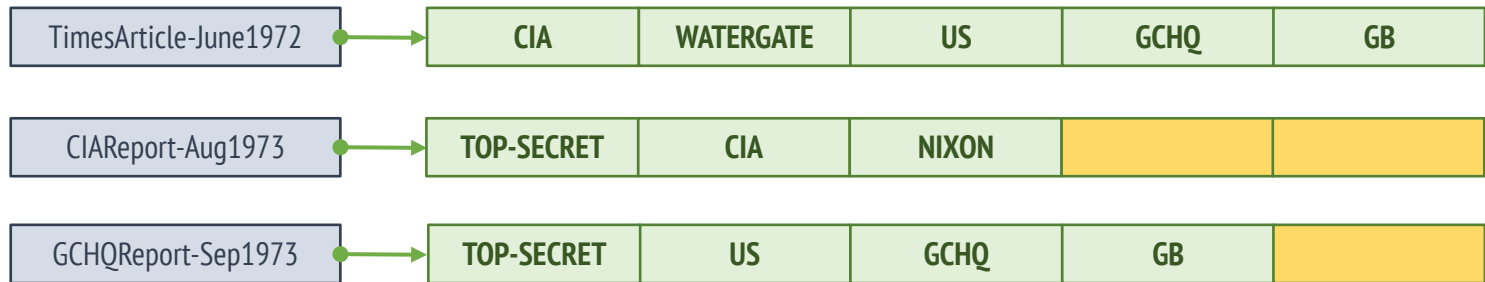
Padding

Forward index



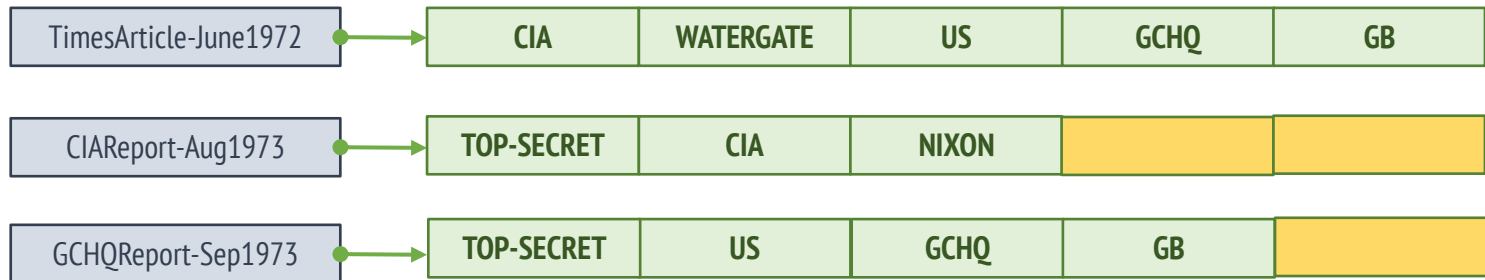
Padding

Forward index

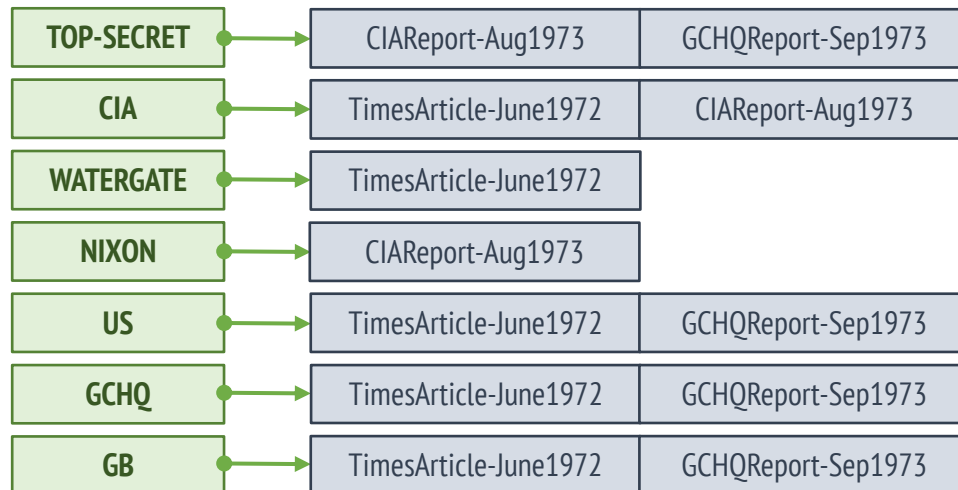


Padding

Forward index

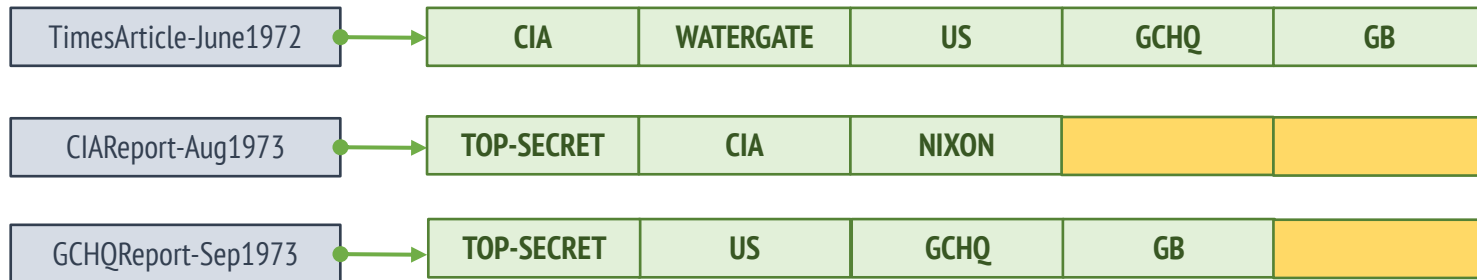


Inverted index

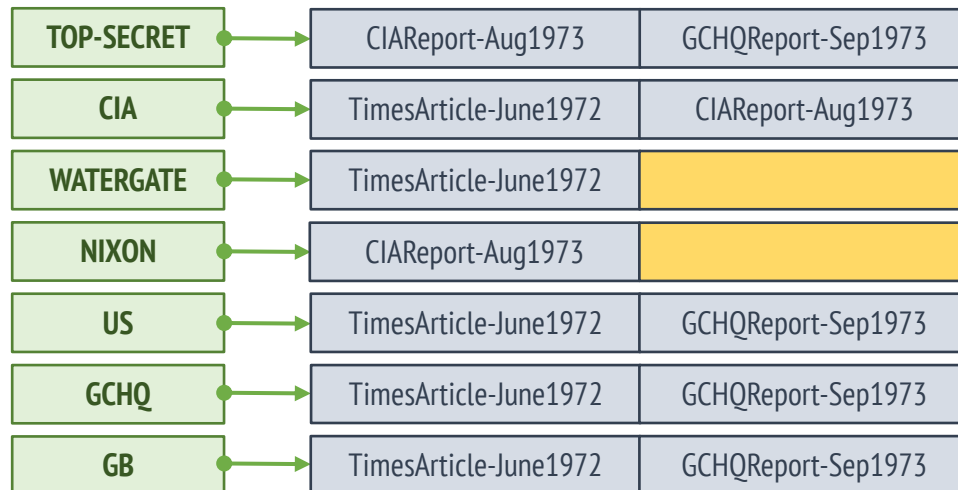


Padding

Forward index

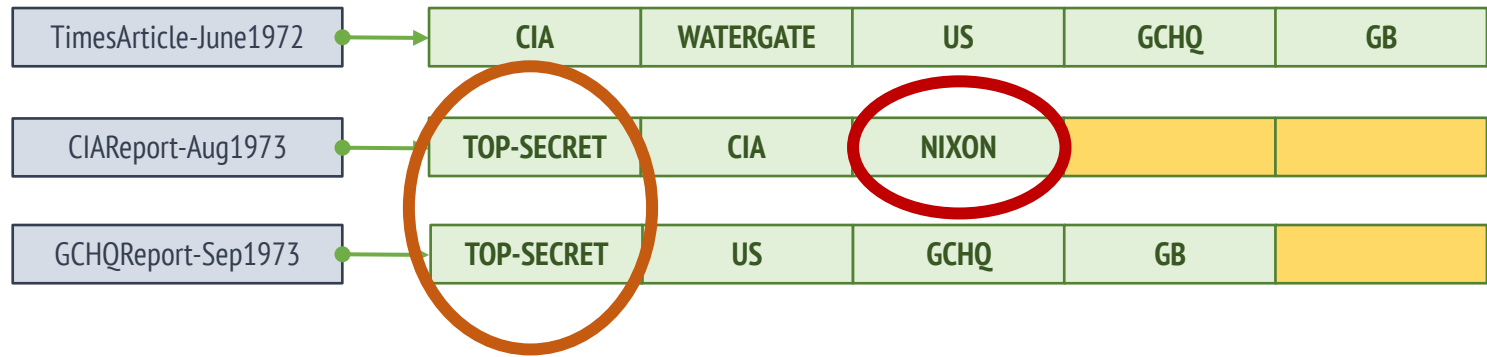


Inverted index



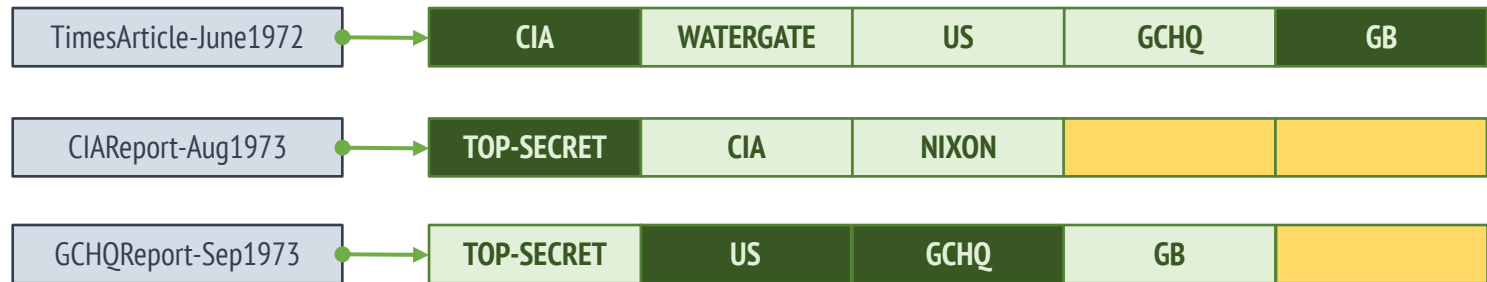
Intersections

Forward index




Intersections

Forward index




Intersections

Forward index



We want



Server the computation

1. Client work needs to be as low as possible.
2. Server needs to do most of the search work.

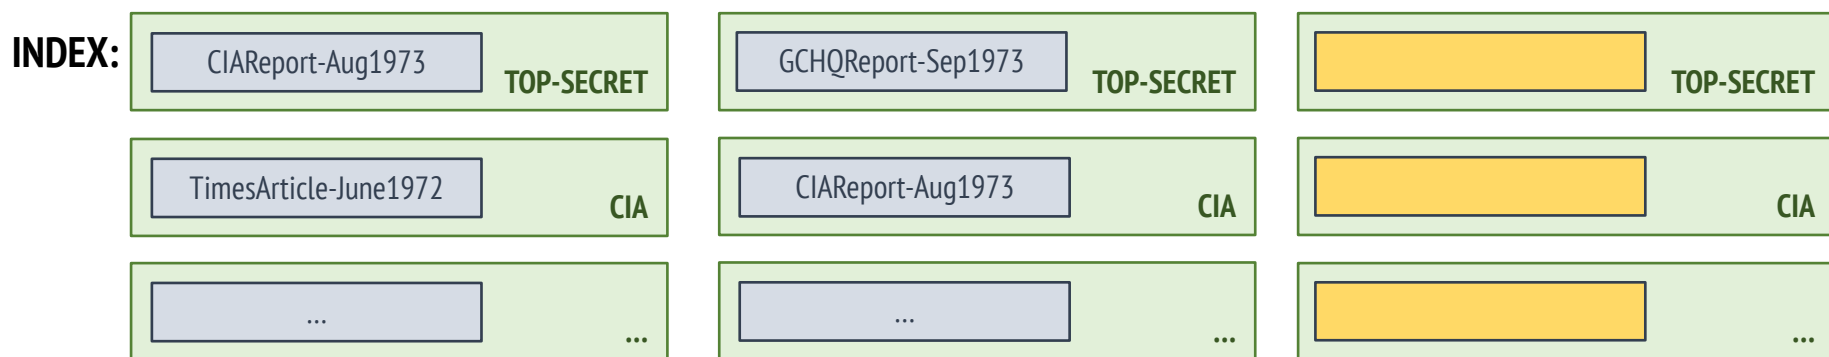
Secure searching

Inverted index:



Symmetric key searchable encryption index:

ENC. DOC. NAMES: CIAReport-Aug1973 GCHQReport-Sep1973 TimesArticle-June1972



Secure searching

Server has:

ENC. DOC. NAMES:

CIAReport-Aug1973

GCHQReport-Sep1973

TimesArticle-June1972

INDEX:

CIAReport-Aug1973

TOP-SECRET

TimesArticle-June1972

CIA

GCHQReport-Sep1973

TOP-SECRET




TOP-SECRET

CIAReport-Aug1973

CIA



CIA



...



...



...

Search term:

TOP-SECRET

Secure searching

Server has:

ENC. DOC. NAMES:

CIAReport-Aug1973

GCHQReport-Sep1973

TimesArticle-June1972

INDEX:

CIAReport-Aug1973

TOP-SECRET

TimesArticle-June1972

CIA

GCHQReport-Sep1973


TOP-SECRET



TOP-SECRET

CIAReport-Aug1973

CIA




CIA



...

...



...

...



...

Search term:

TOP-SECRET

Server computation:

CIAReport-Aug1973

GCHQReport-Sep1973

TimesArticle-June1972

Secure searching

Server has:

ENC. DOC. NAMES:

CIAReport-Aug1973

GCHQReport-Sep1973

TimesArticle-June1972

INDEX:

CIAReport-Aug1973


TOP-SECRET

TimesArticle-June1972

CIA

GCHQReport-Sep1973


TOP-SECRET



TOP-SECRET

CIAReport-Aug1973

CIA



CIA

...

...

...

...

...

Search term:

TOP-SECRET

Server computation:

CIAReport-Aug1973

TOP-SECRET

GCHQReport-Sep1973

TOP-SECRET

TimesArticle-June1972

TOP-SECRET

Secure searching

Server has:

ENC. DOC. NAMES:

CIAReport-Aug1973

GCHQReport-Sep1973

TimesArticle-June1972

INDEX:

CIAReport-Aug1973

TOP-SECRET

TimesArticle-June1972

CIA

GCHQReport-Sep1973

TOP-SECRET



TOP-SECRET

CIAReport-Aug1973


CIA



CIA



...



...



...

Search term:

TOP-SECRET

Server computation:

CIAReport-Aug1973

TOP-SECRET

GCHQReport-Sep1973

TOP-SECRET

TimesArticle-June1972

TOP-SECRET



Secure searching

Server has:

ENC. DOC. NAMES:

CIAReport-Aug1973

GCHQReport-Sep1973

TimesArticle-June1972

INDEX:

CIAReport-Aug1973


TOP-SECRET

TimesArticle-June1972

CIA

GCHQReport-Sep1973


TOP-SECRET




TOP-SECRET

CIAReport-Aug1973

CIA



CIA



...



...



...

Search term:

TOP-SECRET

Server computation:

CIAReport-Aug1973

TOP-SECRET

GCHQReport-Sep1973

TOP-SECRET

TimesArticle-June1972

TOP-SECRET



Result:

CIAReport-Aug1973

GCHQReport-Sep1973

Performance

Example 1 - OXT:

[Cash-Jarecki-Jutla-Krawczyk-Rosu-Steiner13]

- Encrypted database size: 13GB
- DB Contents: 1.5 million emails & attachments
- **Avg. search time: less than 500ms**

Example 2 – 2Lev:

[Cash-Jaeger-Jarecki-Jutla-Krawczyk-Steiner-Rosu14]

- Encrypted database size: 900GBs
- Setup time: 16 hours
- **Avg. query time: less than 200ms**

Searchable encryption limitations

- Encrypted search term is deterministic
- Only search patterns are hidden
- Setting up the index requires a significant amount of time
- Most schemes do not support index extensions

DISCLAIMER :-)

- Searchable encryption **solves** problems related to the security of the search index.
- Searchable encryption **does not solve** problems related to the security of subsequent data retrieval.

Even though the response to the search query has been done in a privacy preserving manner, the server can still learn what the result of the query was by simply observing what the client does next, e.g. **monitor the emails the client is going to access/download.**

Oblivious RAM

Oblivious RAM (ORAM)

- A cryptographic primitive originally designed to prevent reverse engineering by hiding access to memory.
- It has since repurposed for use in client-server scenarios with the purpose of hiding the ways in which data is accessed from the server.

ORAM security requirements

Hide **DATA CONTENTS** and:

1. Hide **which** data is accessed (e.g. My DSS course)
2. Hide **when** data was last accessed (e.g. 5mins ago)
3. Hide **how** data is accessed (i.e. read or write access)
4. Hide **how frequently** data is accessed (e.g. every day at 12pm)
5. Hide the **relationship between consecutive accesses** (e.g. related, random)

ORAM

- Data is stored in blocks of fixed size.
- Uses symmetric encryption (e.g. AES) to encrypt small data structures (e.g. data 'buckets').
- Replaces read and write operations (i.e. download and upload) with a generic **access** operation which contains both a read and a write operation.
- The **access operation** has a significant overhead in order to disguise the exact data being accessed.

ORAM components




- Client stores an AES key
- Client stores a map
- Client stores a stash, i.e., a local cache structure




- Server storage is structured as a binary tree.
- On disk the tree is stored as a flat data structure

Flat binary tree?



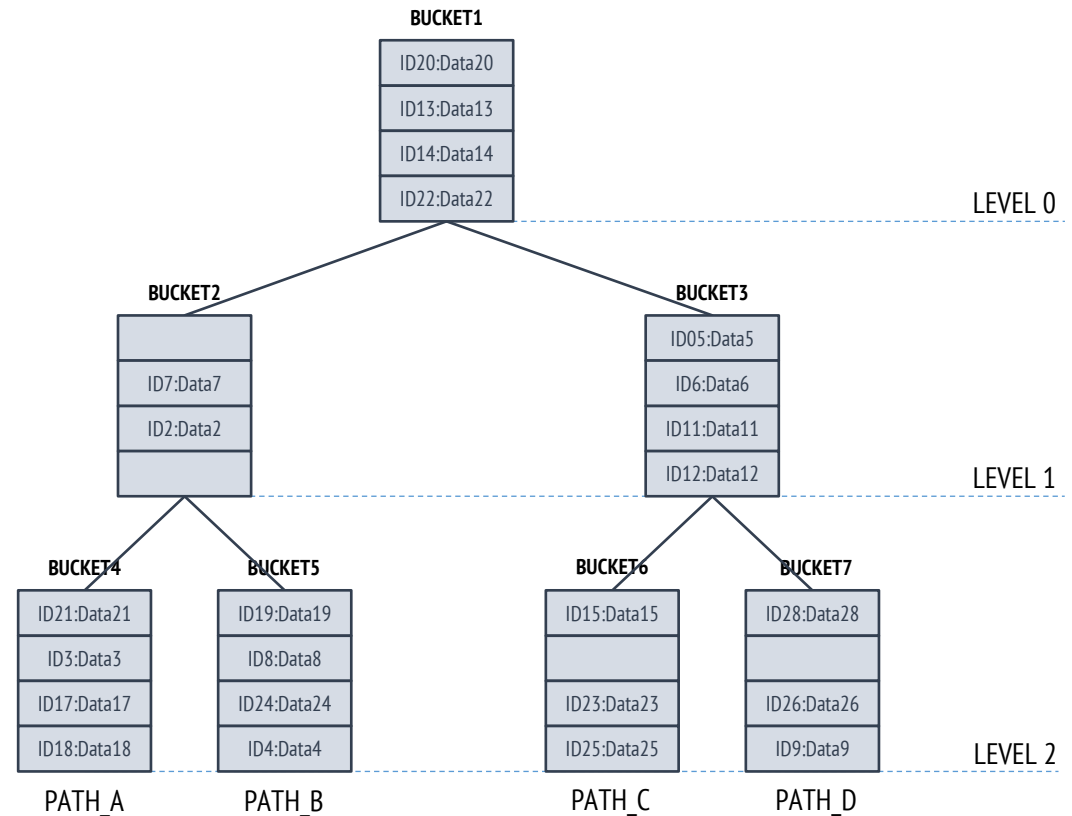

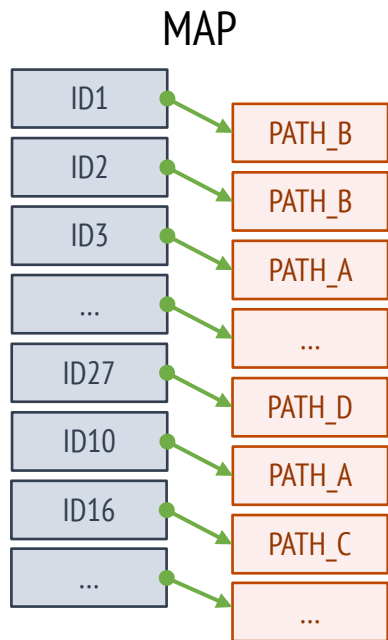
Binary tree



Flat store

PathORAM

[Stefanov-van Dijk-Shi-Chan-Fletcher-Ren-Yu-Devadas13]



PathORAM access

READ:

ID2:empty

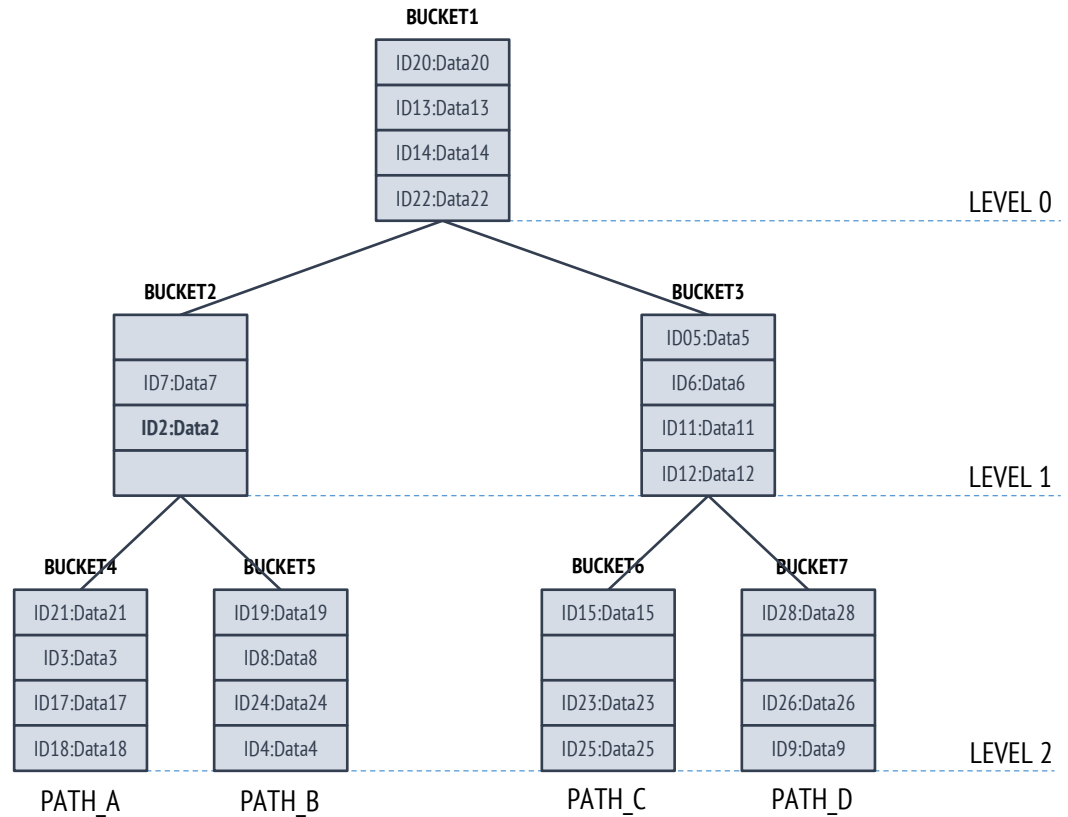
MAP

ID2

PATH_B

REQUEST:

PATH_B



PathORAM access

READ:

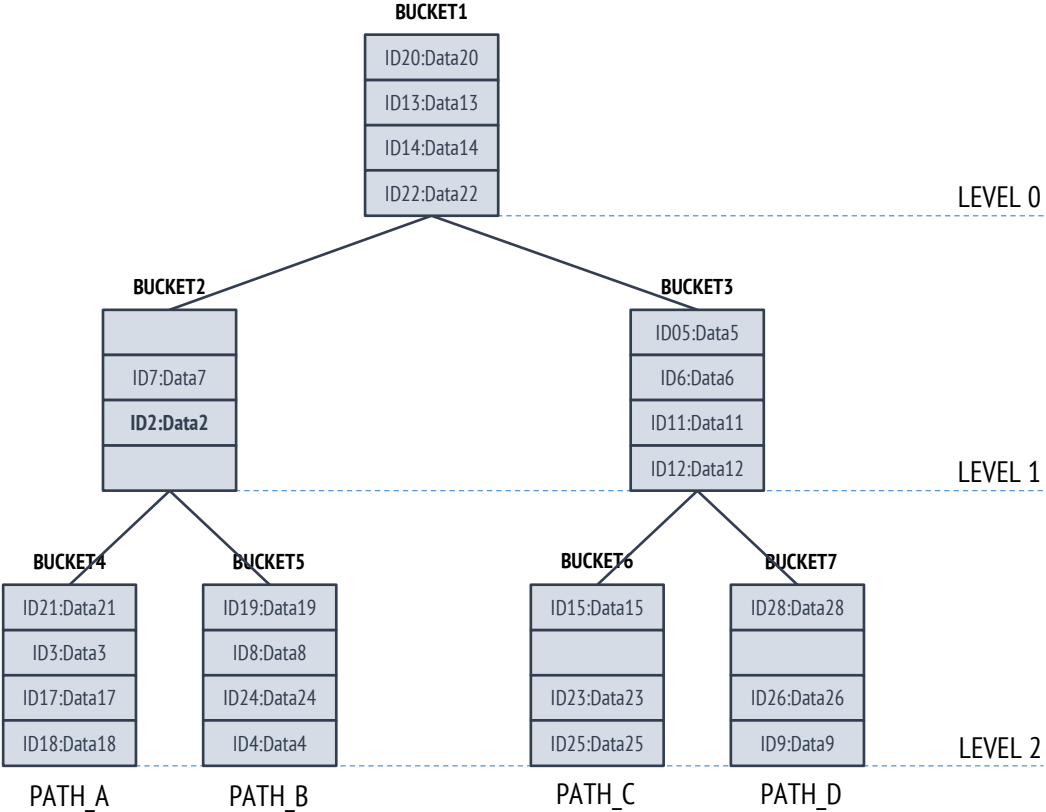
ID2:empty

MAP



REQUEST:

PATH_B



RECEIVE:

PATH_B											
Bucket5				Bucket2				Bucket1			
ID4:Data4	ID24:Data24	ID8:Data8	ID19:Data19		ID2:Data2	ID7:Data7		ID22:Data22	ID14:Data14	ID13:Data13	ID20:Data20

PathORAM access

READ:

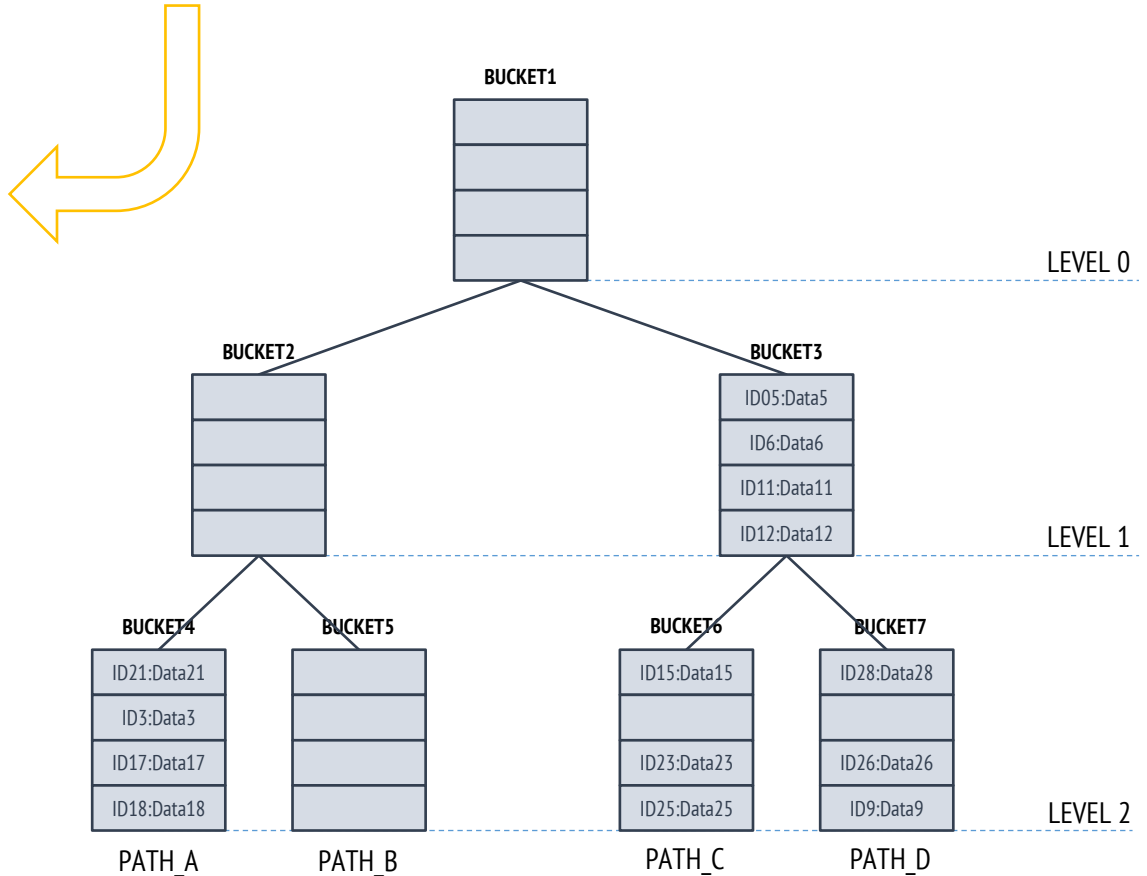
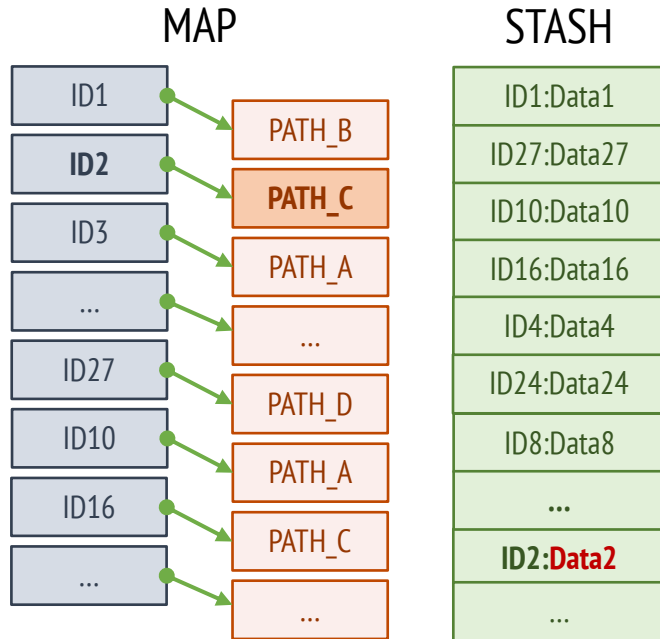
ID2:empty

REQUEST:

PATH_B

RECEIVE:

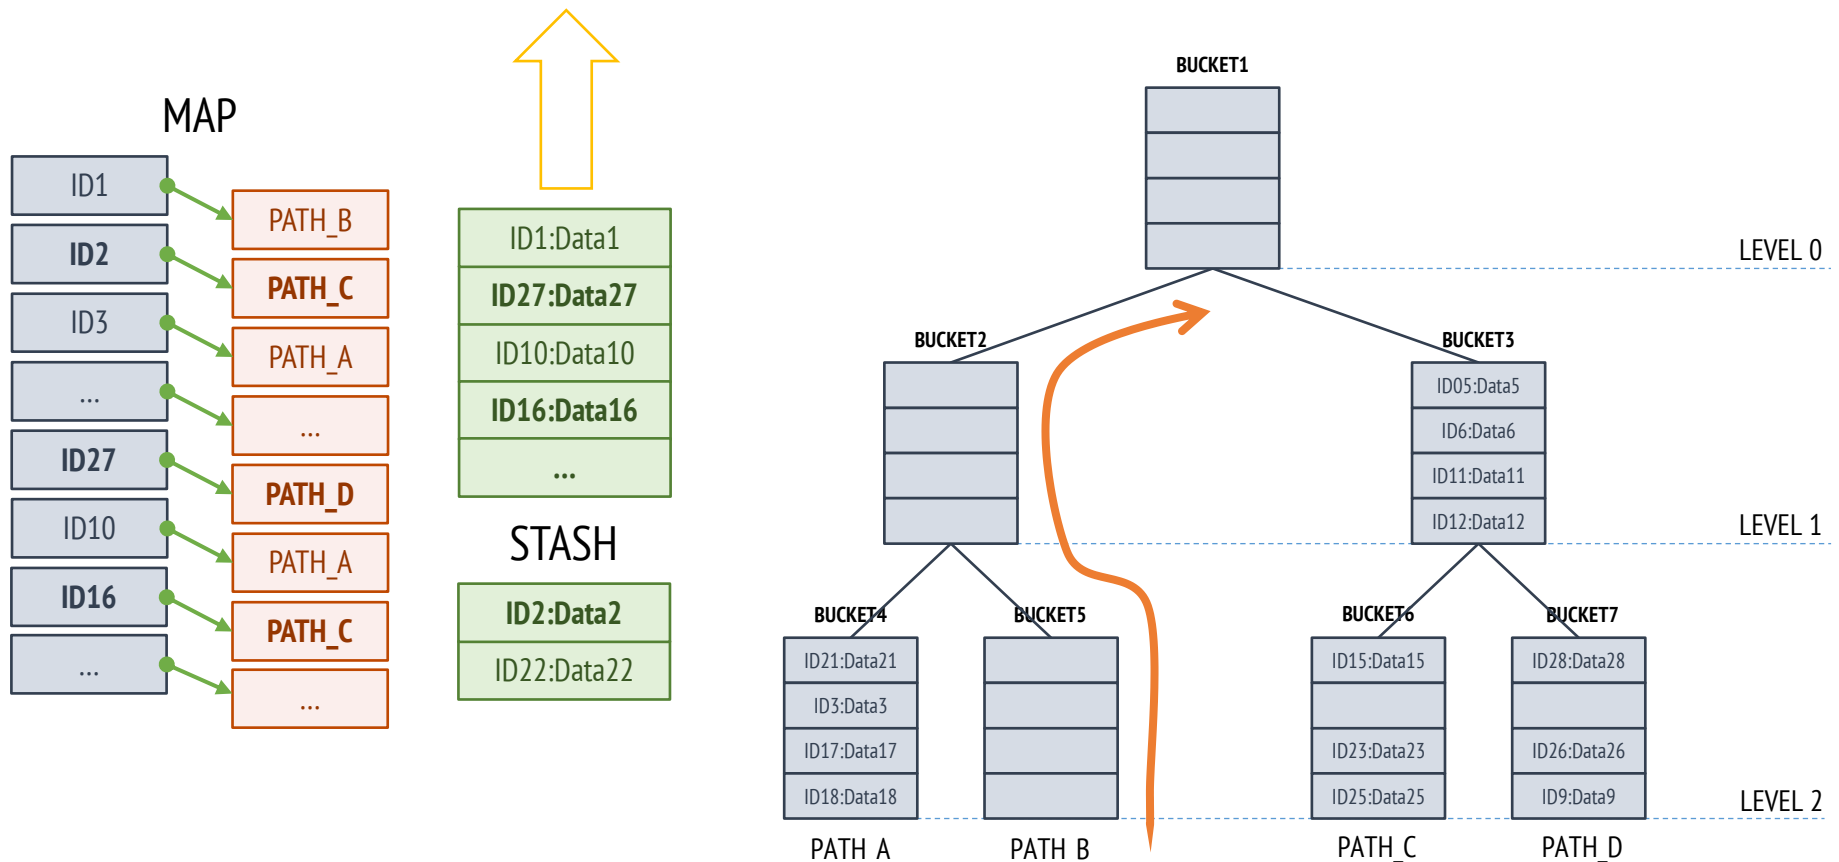
PATH_B											
Bucket5				Bucket2				Bucket1			
ID4:Data4	ID24:Data24	ID8:Data8	ID19:Data19		ID2:Data2	ID7:Data7		ID22:Data22	ID14:Data14	ID13:Data13	ID20:Data20



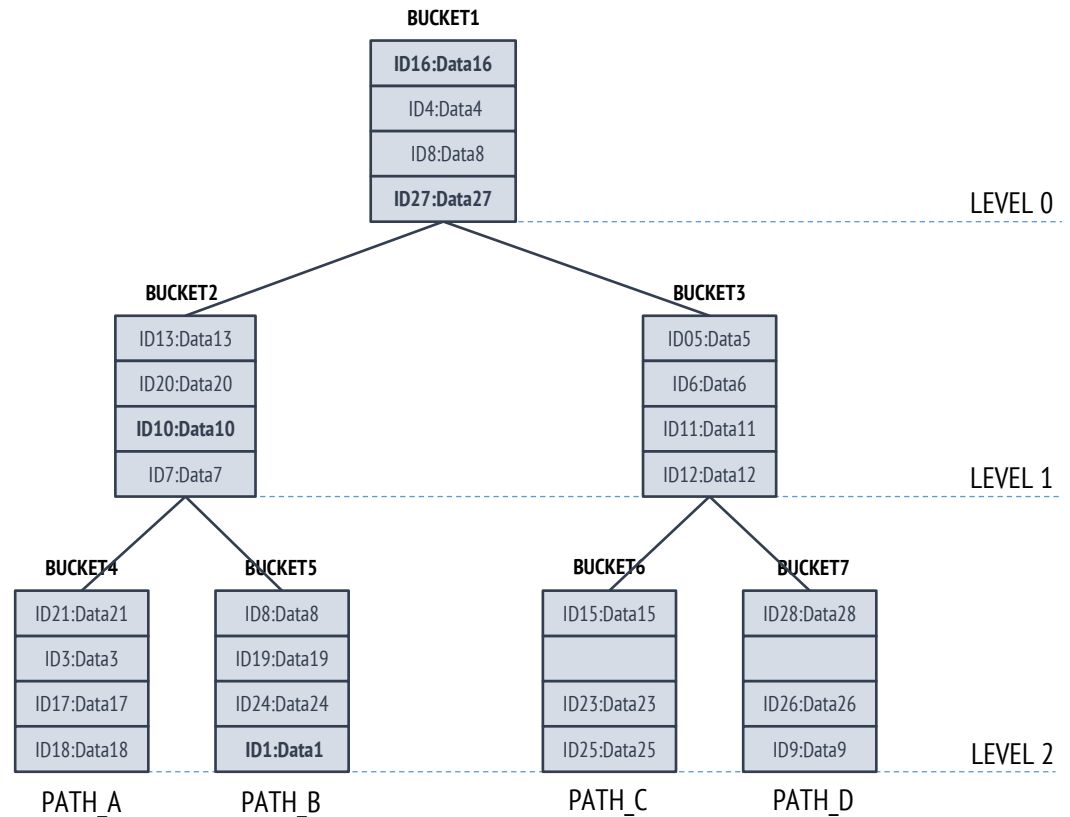

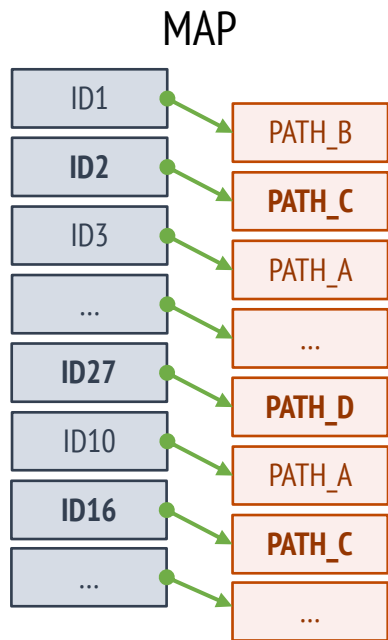
PathORAM access

WRITE:

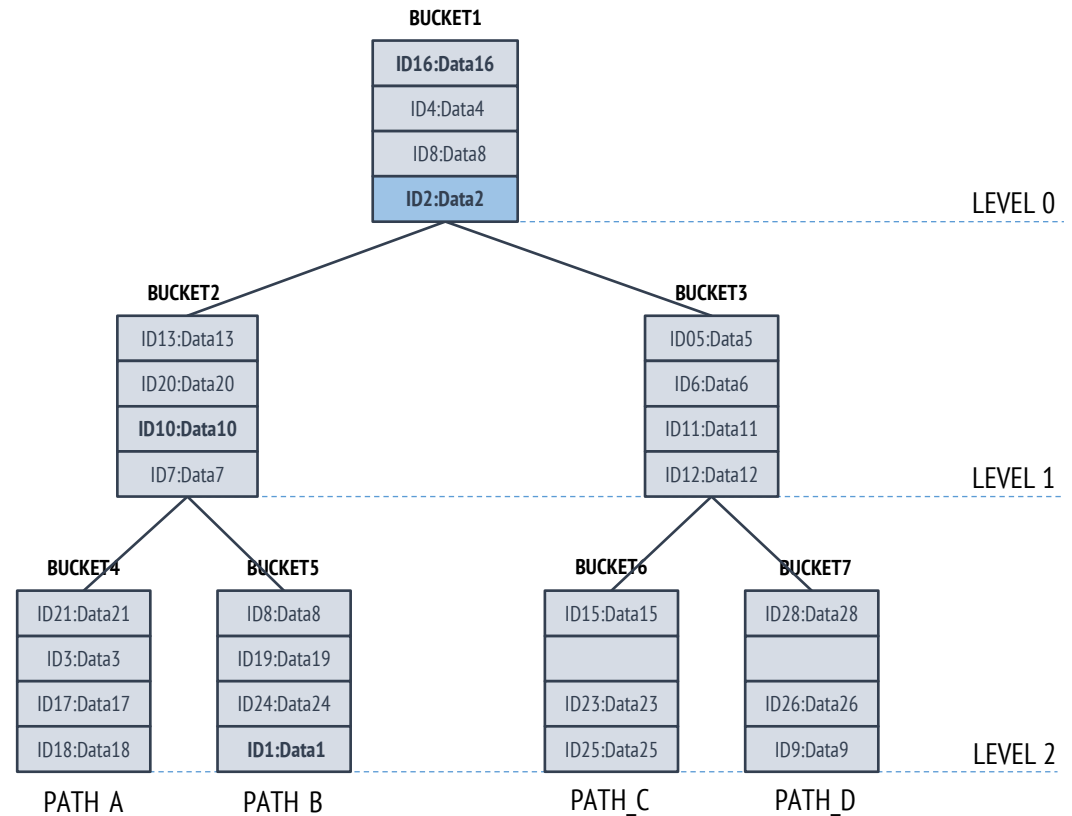


PATH_B											
Bucket5				Bucket2				Bucket1			
ID1:Data1	ID24:Data24	ID19:Data19	ID8:Data8	ID7:Data7	ID10:Data10	ID20:Data20	ID13:Data13	ID27:Data27	ID8:Data8	ID4:Data4	ID16:Data16



PathORAM structure



PathORAM structure (alternative)



Algorithm




Access(op, a, data^{*}):

- 1: $x \leftarrow \text{position}[a]$
- 2: $\text{position}[a] \leftarrow \text{UniformRandom}(0 \dots 2^L - 1)$
- 3: **for** $\ell \in \{0, 1, \dots, L\}$ **do**
- 4: $S \leftarrow S \cup \text{ReadBucket}(\mathcal{P}(x, \ell))$
- 5: **end for**
- 6: $\text{data} \leftarrow \text{Read block } a \text{ from } S$
- 7: **if** $\text{op} = \text{write}$ **then**
- 8: $S \leftarrow (S - \{(a, \text{data})\}) \cup \{(a, \text{data}^*)\}$
- 9: **end if**
- 10: **for** $\ell \in \{L, L - 1, \dots, 0\}$ **do**
- 11: $S' \leftarrow \{(a', \text{data}') \in S : \mathcal{P}(x, \ell) = \mathcal{P}(\text{position}[a'], \ell)\}$
- 12: $S' \leftarrow \text{Select } \min(|S'|, Z) \text{ blocks from } S'.$
- 13: $S \leftarrow S - S'$
- 14: $\text{WriteBucket}(\mathcal{P}(x, \ell), S')$
- 15: **end for**
- 16: **return** data


Some shortcomings

- Client has to store a map
 - **Can be potentially big!**
- Client has to store a stash
 - Size of a a full path ($O(\log N)$ complexity)
 - PathORAM uses an aggressive stash emptying strategy

Recursive ORAM



Malicious adversary ORAM (using a modified Merkle tree)



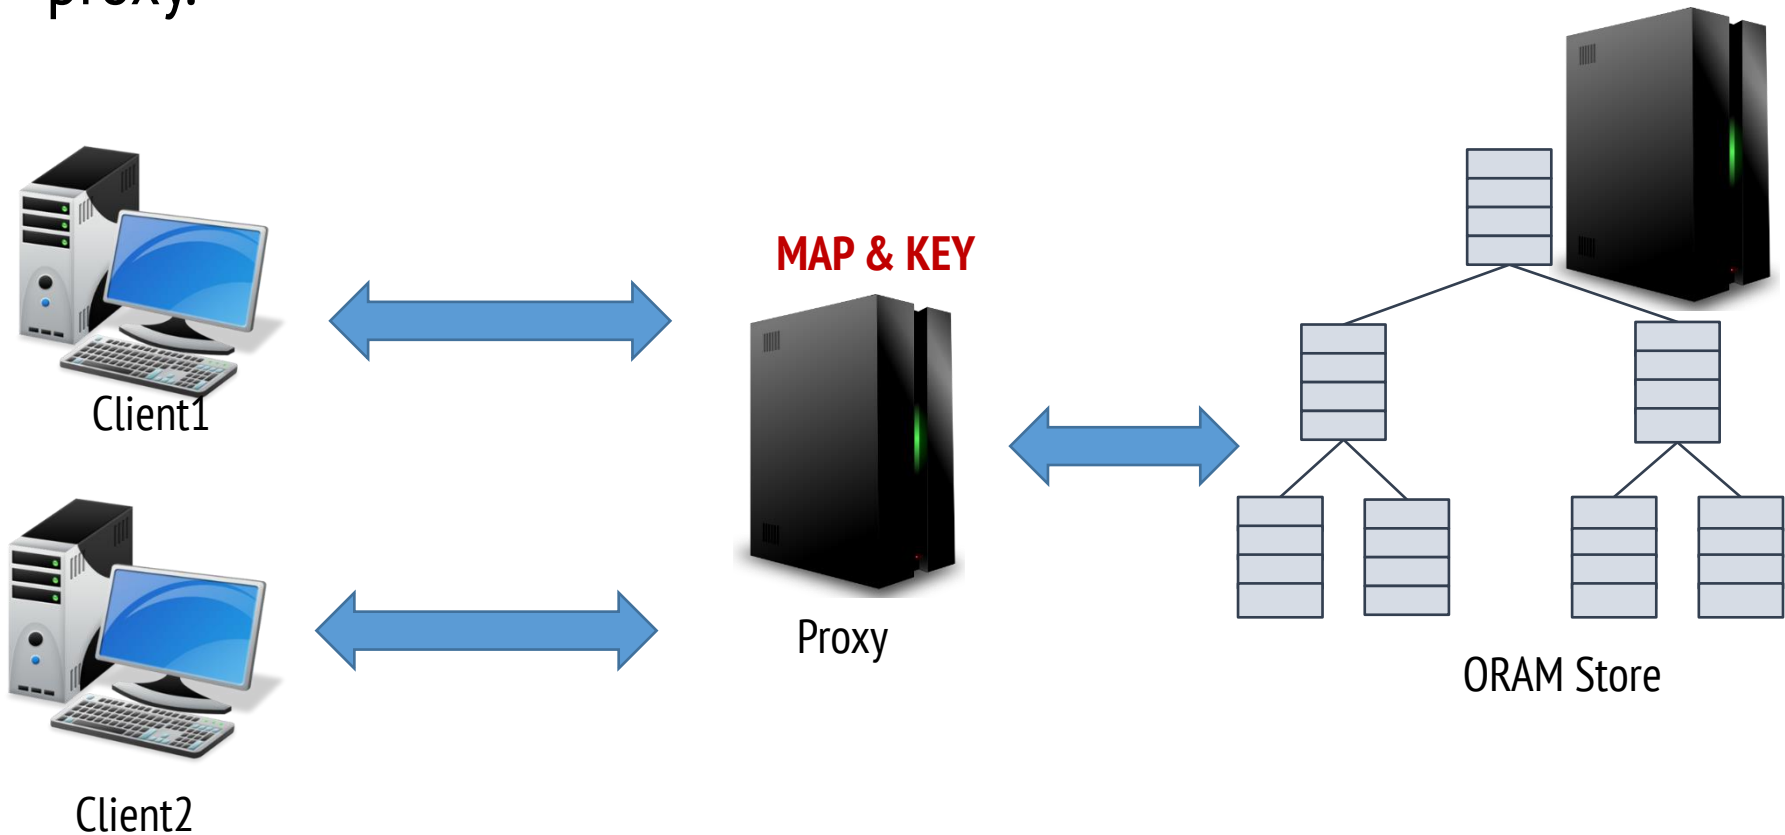
Integrity records:

$$H_N = H(b_1 \dots b_n, H_{C1}, H_{C2})$$

$$H_L = H(b_1 \dots b_n, 0, 0,)$$

Other Limitations

PathORAM is limited to a single user. If multiple users require access to the store (server), access must be done through a proxy.



Other Limitations

If multiple users access the store timing attacks can be leveraged by the server with respect to

1. Proxy data CACHING
2. Proxy duplicating requests (e.g. Client1 and Client2 request same data)
3. Volume of data (e.g. Client1 wants more data than Client2)

PathORAM performance

Example

Assuming a 128GB database with:

- $S = 64\text{KB}$ block size
- $Z = 5$ blocks per bucket
- $L = 20$ levels

`SecretDocument.txt`


a 1MB document stored in the database

PathORAM performance

Example

Assuming a 128GB database with:

- $S = 64\text{KB}$ block size
- $Z = 5$ blocks per bucket
- $L = 20$ levels



What are the bandwidth requirements to access this document?

`SecretDocument.txt`


a 1MB document stored in the database

PathORAM performance

Example

Assuming a 128GB database with:

- $S = 64\text{KB}$ block size
- $Z = 5$ blocks per bucket
- $L = 20$ levels



$1\text{MB} = 1024\text{KB}$

Block per document N :

$N = 1024\text{KB} / 64\text{KB} \text{ (size of the block)} = 16$

SecretDocument.txt


a 1MB document stored in the database

PathORAM performance

Example

Assuming a 128GB database with:

- $S = 64\text{KB}$ block size
- $Z = 5$ blocks per bucket
- $L = 20$ levels



$1\text{MB} = 1024\text{KB}$

Block per document N :

$N = 1024\text{KB} / 64\text{KB}$ (size of the block) = 16

To send/receive ONE document

PathORAM requires: $N * S * Z * L = 100\text{MB}$

`SecretDocument.txt`

a 1MB document stored in the database

ORAM applications

- Personal health records
- Credit score systems
- GENOME related research
- As a private information retrieval (PIR) protocol

ORAM vs. Searchable Encryption

ORAM

- Provides anonymous access to data blocks
- Used as a private information retrieval (PIR) protocol
- Fully protects access patterns and data contents
- Requires a considerable overheads

Searchable encryption

- Enables users to securely search a precomputed index
- Used to efficiently **locate** data in databases
- Protects search terms and search results
- Only protects search patterns