


Protection and Security

- Two main features on computer processors allow operating systems to provide protection and security
- Feature 1: **multiple processor operating modes**
 - The processor physically enforces different constraints on programs operating in different modes
- Minimal requirements:
 - Kernel mode (a.k.a. protected mode, privileged mode, etc.) allows a program full access to all processor capabilities and operations
 - User mode (a.k.a. normal mode) only allows a program to use a restricted subset of processor capabilities
- The kernel of the operating system is the part of the OS that runs in kernel mode
 - The OS may have [many] other components running in user mode

Protection and Security

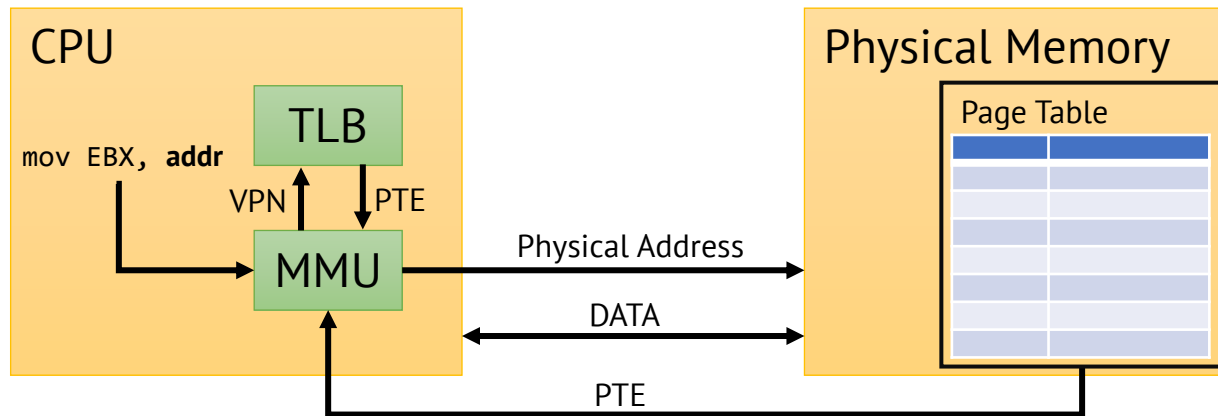
- Some processors provide more than two operating modes
- These are called hierarchical protection domains or protection rings:
 - Higher-privilege rings can also access lower-privilege operations and data
- x86 provides four operating modes
 - Level 0 = kernel mode; level 3 = user mode
- Support for multiple protection levels is ubiquitous, even in mobile devices
 - e.g. ARMv7 processors in modern smartphones have 8 different protection levels for different scenarios



Protection and Security

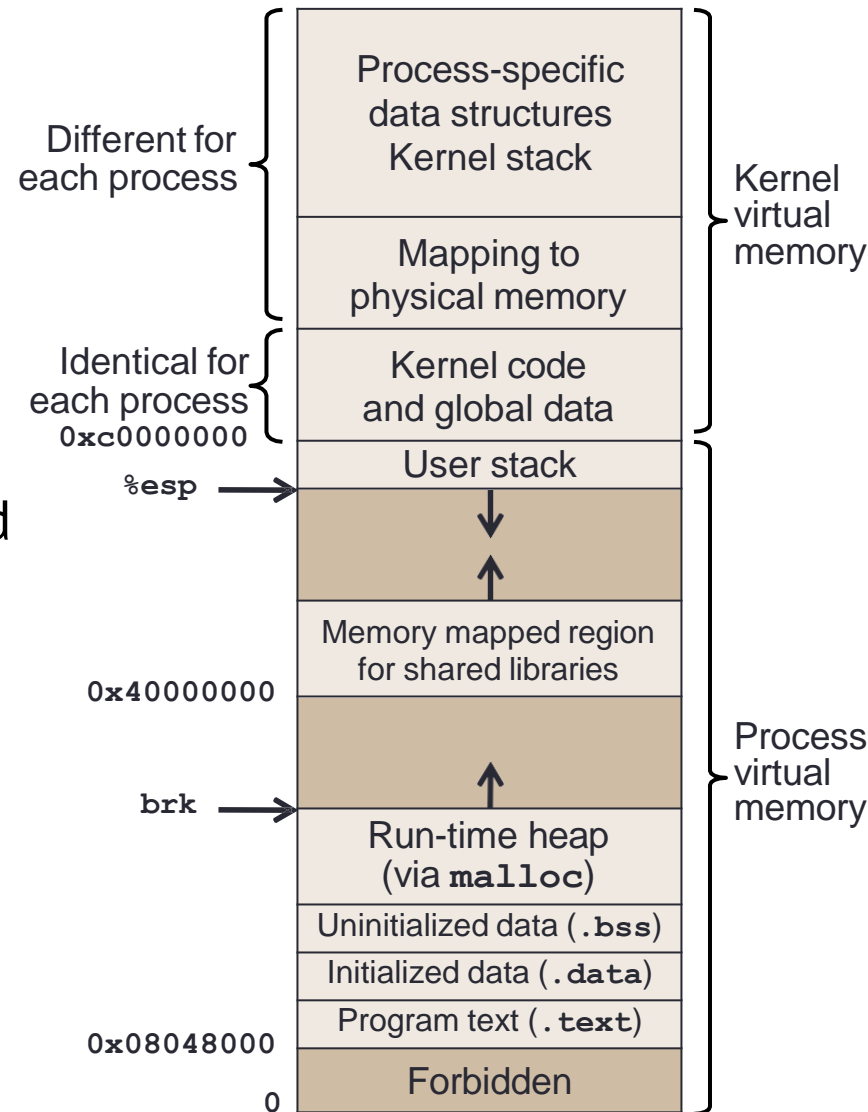
- Feature 2: **virtual memory**

- The processor maps virtual addresses to physical addresses using a page table
- The Memory Management Unit (MMU) performs this translation
- Translation Lookaside Buffers (TLBs) cache page table entries to avoid memory access overhead when translating addresses
- Only the kernel can manipulate the MMU's configuration.



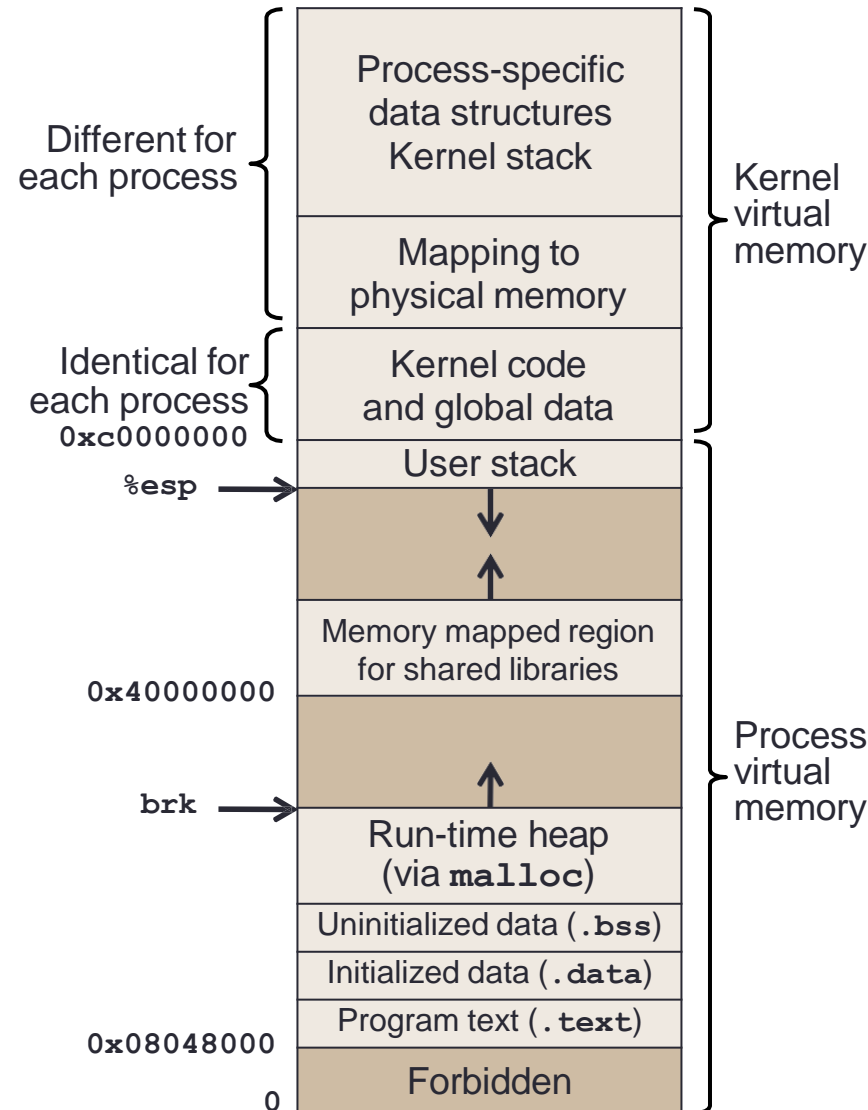
Protection and Security

- Virtual memory allows OS to give each process its own isolated address space
 - Processes have identical memory layouts, simplifying compilation, linking and loading
- Regions of memory can also be restricted to kernel-mode access only, or allow user-mode access
 - Called kernel space and user space
 - If user-mode code tries to access kernel space, processor notifies the OS
 - Only kernel can manipulate this config!



Protection and Security

- The OS must track certain details for each process
 - e.g. process' memory mapping
 - e.g. the process' scheduling configuration and behavior
- A process can't be allowed to access these details directly!
 - Just as with global kernel state, allowing direct access would open security holes
 - Processes must ask the kernel to manipulate these states on their behalf



Example: Console and File I/O

- You run a program on a Windows or UNIX system
 - The OS sets up certain basic facilities for your program to use
- Standard input/output/error streams
 - What **printf()** and **scanf()** use by default!
- Standard input/output/error streams can be from:
 - The console/terminal
 - Redirected to/from disk files
 - Your program sees the contents of a disk file on its standard input
 - What your program writes on standard output goes to a file on disk
- Redirected to/from another process!
 - Your program sees output of another process on its standard input
 - Your program's standard output is fed to another process' standard input

UNIX File IO


- All input/output is performed with UNIX system functions:

```
ssize_t read(int filedes, void *buf, size_t nbyte)  
ssize_t write(int filedes, const void *buf, size_t nbyte)
```

- Attempt to read or write nbyte bytes to file specified by **filedes**
 - Actual number of bytes read or written is returned by the function
 - EOF indicated by 0 return-value; errors indicated by values < 0
- The user program requests that the kernel reads or writes up to nbyte bytes, on behalf of the process:
 - **read()** and **write()** are system calls
 - Frequently takes a long time (milliseconds or microseconds; even more for user input)
 - Kernel often initiates the request, then context-switches to another process until I/O subsystem fires an interrupt to signal completion


UNIX File IO

- **filedes** is a file descriptor
 - A non-negative integer value that represents a specific file or device
- Processes can have multiple open files
 - Each process' open files are recorded in an array of pointers
 - Array elements point to file structs describing the open file, e.g. the process' current read/write offset within the file
 - **filedes** is simply an index into this array
 - (Each process has a cap on total # of open files)
- Every process has this data structure, but processes are not allowed to directly manipulate it!
 - The kernel maintains this data structure on behalf of each process



UNIX File IO

- Individual **file** structs reference the actual details of how to interact with the file
 - Allows OS to support many kinds of file objects, not just disk files




file_ops is a struct containing function-pointers for common operations supported by all file types, e.g.

```
struct file_operations {  
    ssize_t (*read)(file *f, void *buf, size_t nb);  
    ssize_t (*write)(file *f, void *buf, size_t nb);  
    ...  
};
```

UNIX File/Console IO

- Individual file structs reference the actual details of how to interact with the file
 - Allows OS to support many kinds of file objects, not just disk files

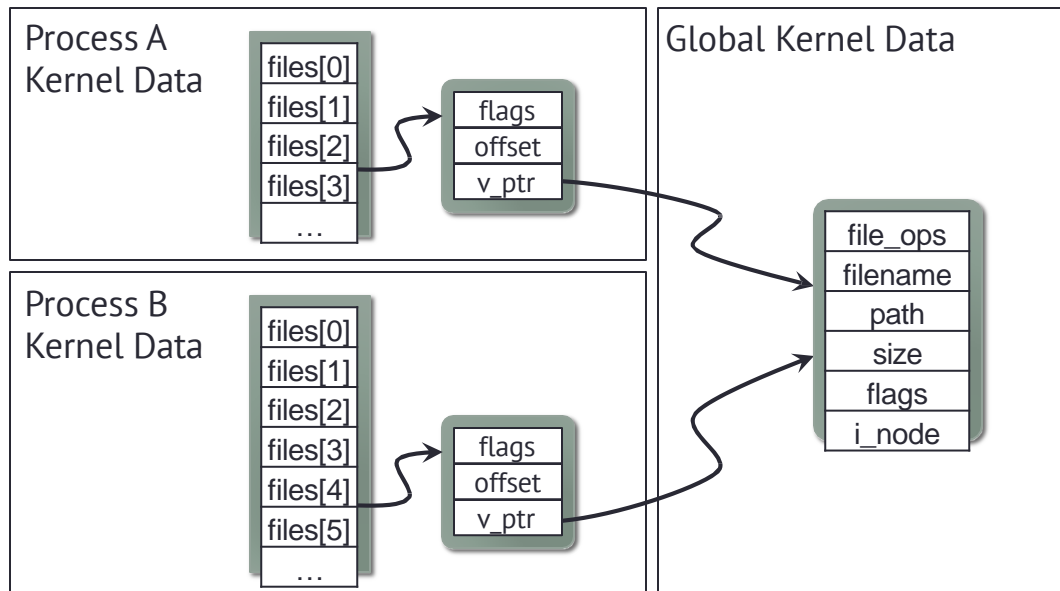


- Kernel can easily read and write completely different file types using indirection

```
// Kernel code for read(filedes,buf,nbyte)
file *f = files[filedes];
f->v_ptr->file_ops->read(file, buf, nbyte);
```

UNIX File IO

- Levels of indirection also allow multiple processes to have the same file open
 - Each process has its own read/write offset for the file
 - Operations are performed against the same underlying disk file



UNIX Standard I/O

- When a UNIX process is initialized by the OS, standard input/output/error streams are set up automatically
- Almost always:
 - File descriptor 0** = standard input
 - File descriptor 1** = standard output
 - File descriptor 2** = standard error
- For sake of compatibility, always use constants defined in **unistd.h**
- standard header file
 - STDIN_FILENO** = file descriptor of standard input
 - STDOUT_FILENO** = file descriptor of standard output
 - STDERR_FILENO** = file descriptor of standard error

UNIX Standard I/O and Command Shells

- Most programs don't really care about where **stdin** and **stdout** go, if they work
 - Command shells care very much!
- **grep Allow < logfile.txt > output.txt**
 - Shell sets grep's **stdin** to read from logfile.txt
 - Shell sets grep's **stdout** to write to the file output.txt
 - (If output.txt exists, it is truncated)
- Once **stdin** and **stdout** are properly set, grep is invoked:
argc = 2, argv = {"grep", "Allow", NULL}